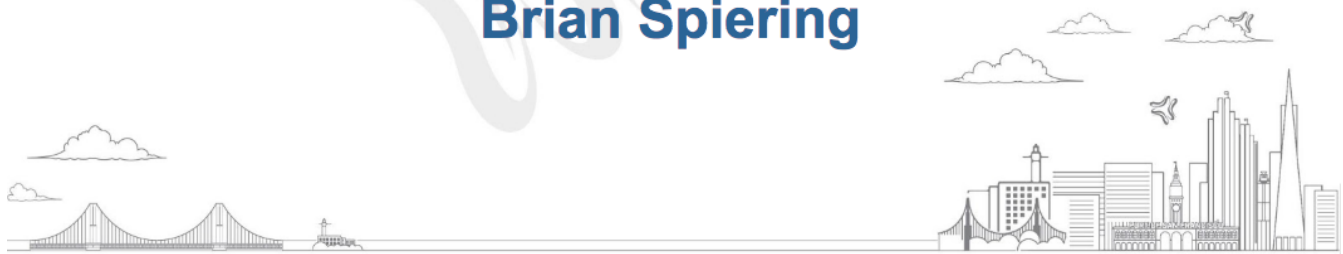




# An Absolute Beginner's Guide to Deep Learning with Keras

**Brian Spiering**



3rd Annual Regional Python Conference

August 16 - 19, 2018 | San Francisco, CA

[bit.ly/pybay-keras](http://bit.ly/pybay-keras) (<http://bit.ly/pybay-keras>)

## What Do I Do?

Professor @



UNIVERSITY OF  
SAN FRANCISCO

Master of Science  
in Data Science

# Deep Learning



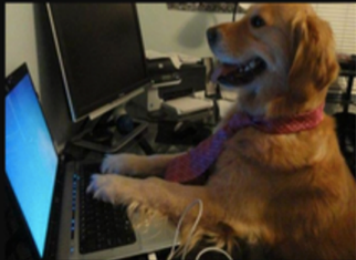
What society thinks I do



What my friends think I do



What other computer scientists think I do



What mathematicians think I do



What I think I do

```
from keras import *
```

What I actually do

## Keras - Neural Networks for humans



A high-level, intuitive API for Deep Learning.

Easy to define neural networks, then automatically handles execution.

A simple, modular interface which allows focus on learning and enables fast experimentation.

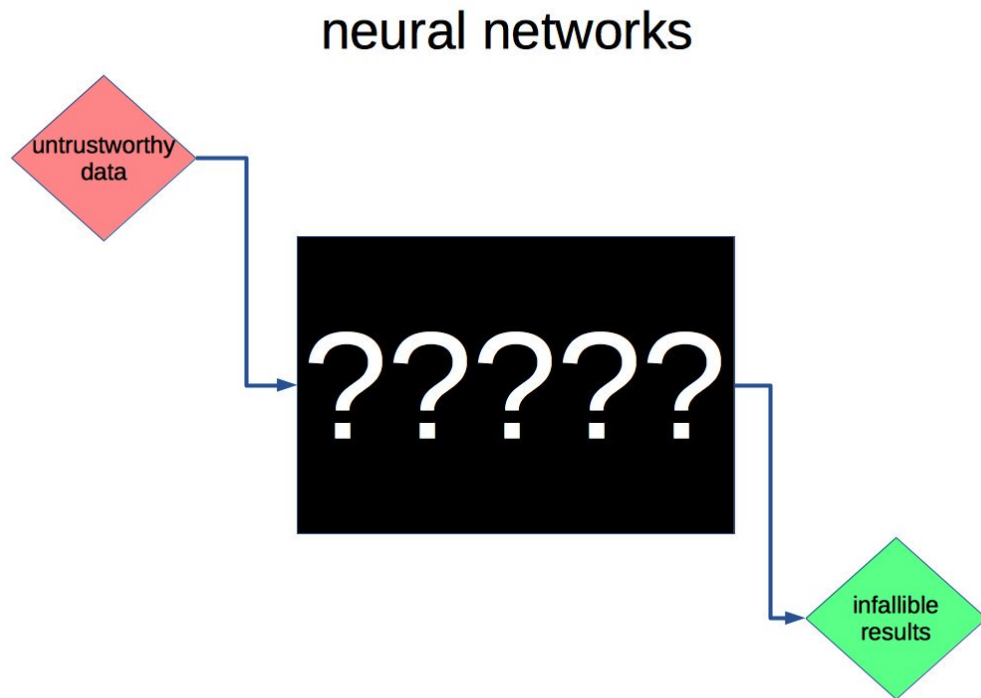
## Goals

- General introduction to Deep Learning
- Overview of Keras library
- An end-to-end example in Keras

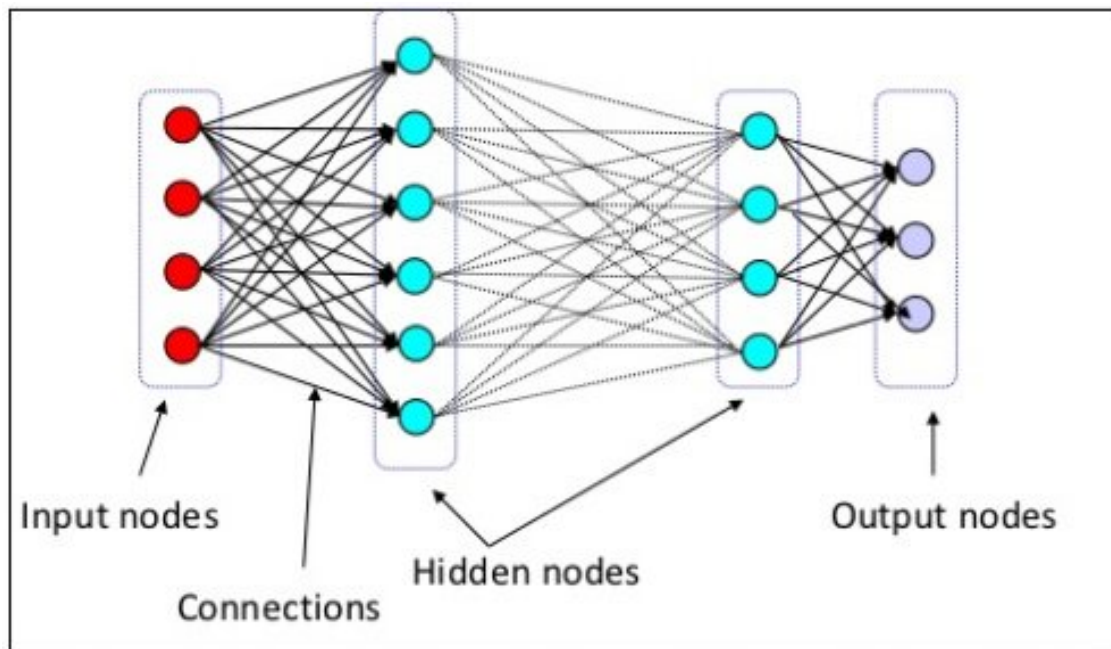
## Anti-Goals

- Understanding of Deep Learning
- Building neural networks from scratch
- A complete survey of keras library

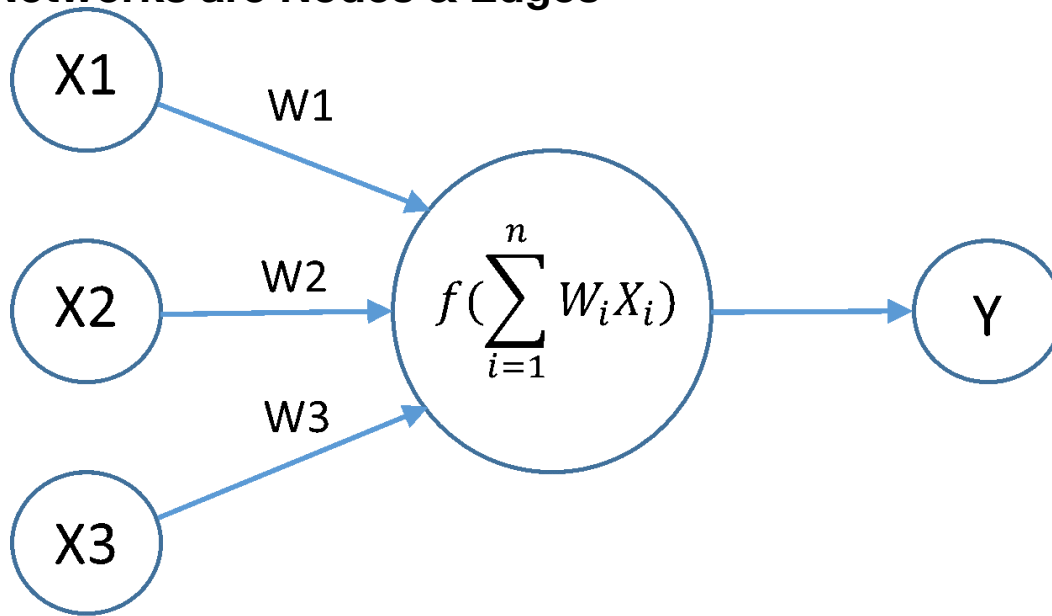
## Deep Learning 101



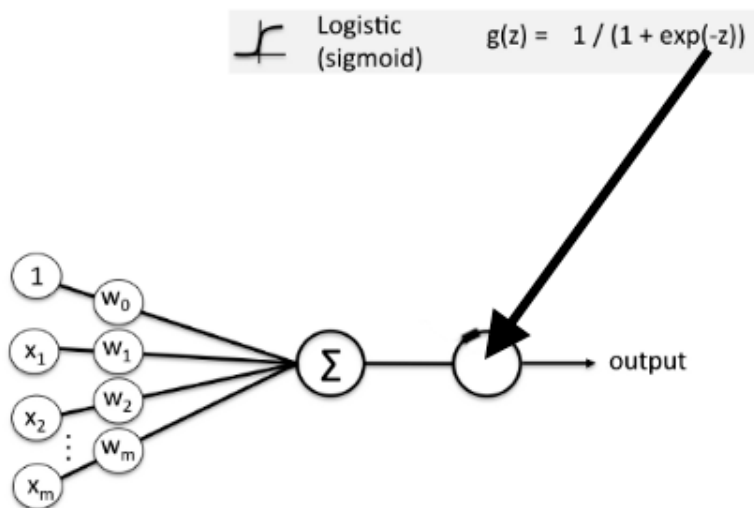
**Deep Learning (DL) are Neural networks (NN) with  $>1$  hidden layer**



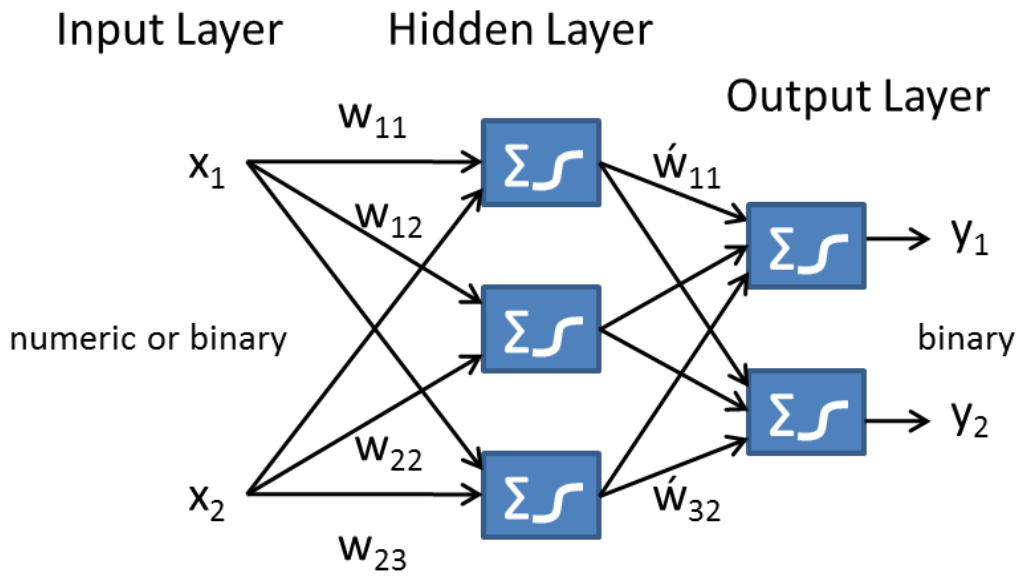
## Neural Networks are Nodes & Edges



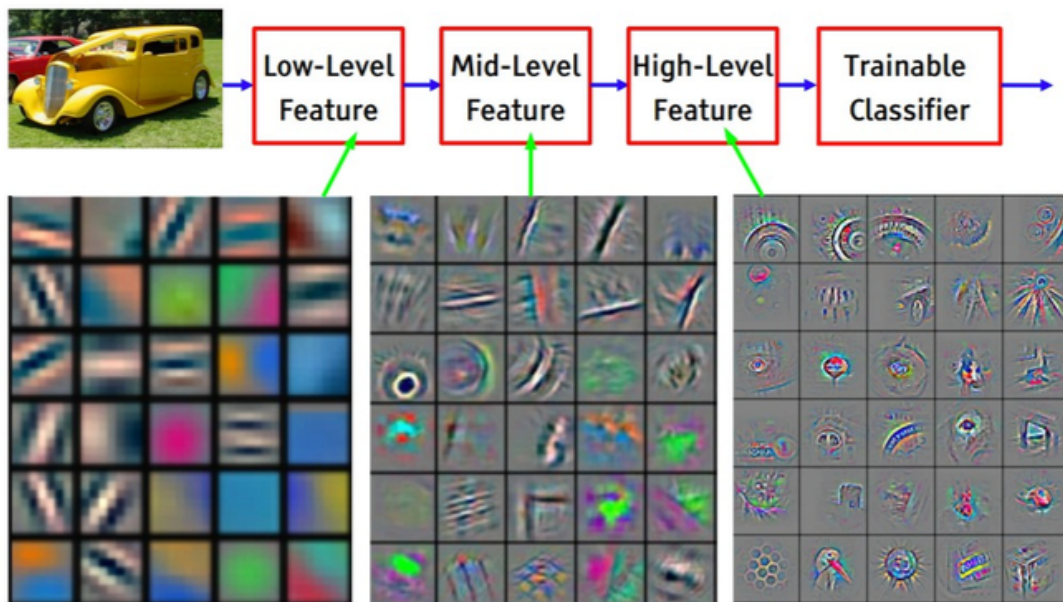
## Nonlinear function allows learning of nonlinear relationships



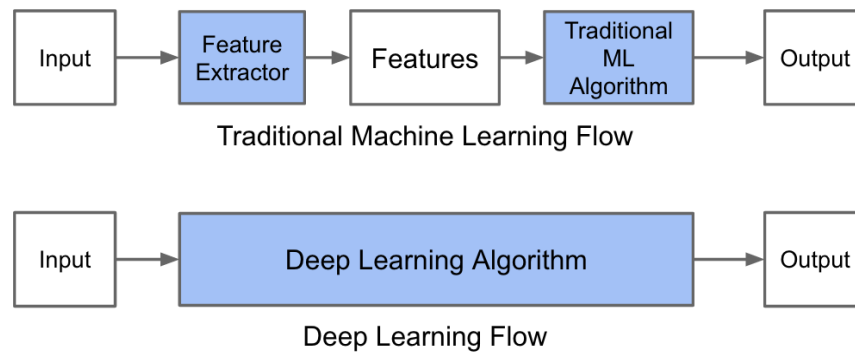
## Groups of nodes all the way down



**Deep Learning isn't magic, it is just very good at finding patterns**



## Deep Learning has fewer steps than traditional Machine Learning



## If you want to follow along...

GitHub repo: [bit.ly/pybay-keras](http://bit.ly/pybay-keras) (<http://bit.ly/pybay-keras>)

## If you want to type along...

1. Run a local Jupyter Notebook
2. [Binder \(https://mybinder.org/v2/gh/brianspiering/keras-intro/master\)](https://mybinder.org/v2/gh/brianspiering/keras-intro/master): In-Browser Jupyter Notebook
3. [Colaboratory \(https://colab.research.google.com/\)](https://colab.research.google.com/): "Google Docs for Jupyter Notebooks"

```
In [339]: reset -fs
```

```
In [340]: import keras
```

```
In [341]: # What is the backend / execution engine?
```

```
In [342]: keras.backend.backend()
```

```
Out[342]: 'tensorflow'
```



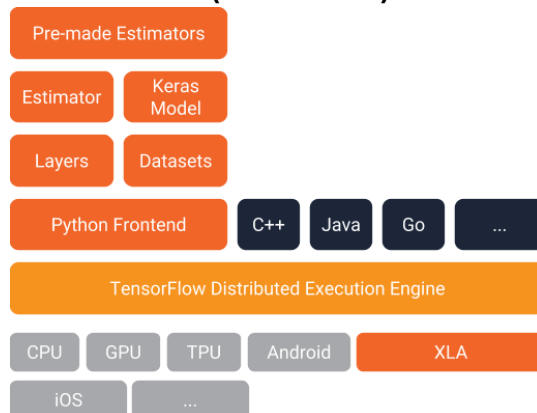
**An open-source software library for Machine Intelligence**

**Numerical computation using data-flow graphs (DFG)**

## TensorFlow: A great backend

**A very flexible architecture which allows you to do almost any numerical operation.**

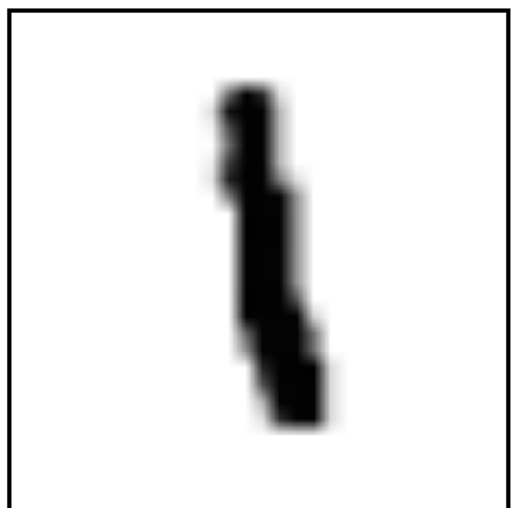
**Then deploy the computation to CPUs or GPUs (one or more) across desktop, cloud, or mobile device.**





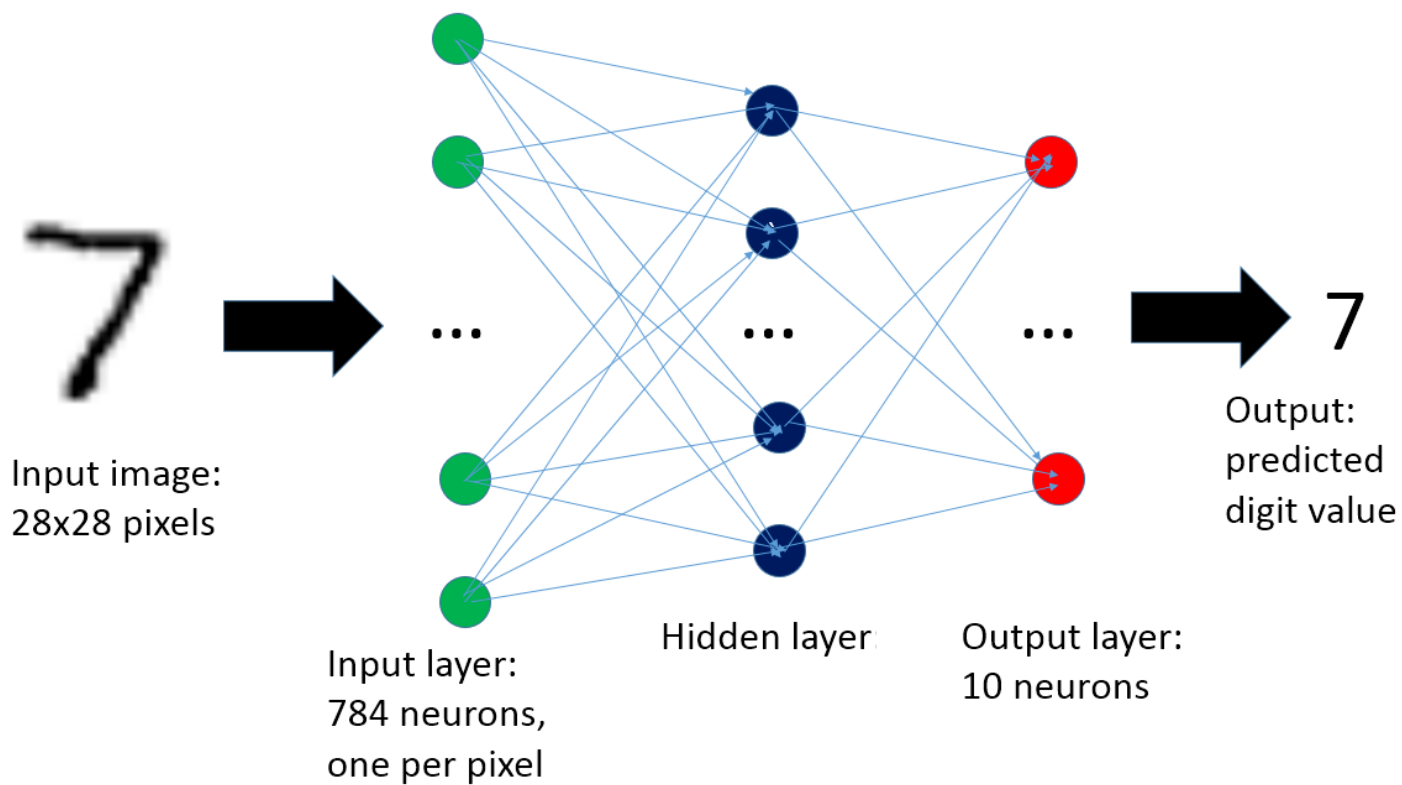
MNIST handwritten digit database:  
The "Hello World!" of Computer Vision

2	1	0	4	1	4	9	5
9	0	6	9	0	1	5	9
7	3	4	9	6	6	5	4
0	7	4	0	1	3	1	3
4	7	2	7	1	2	1	1
7	4	2	3	5	1	2	4
4	6	3	5	5	6	0	4
1	9	5	7	8	9	3	7



21

[illegible]



```
In [343]: # Import data
```

```
In [344]: from keras.datasets import mnist
```

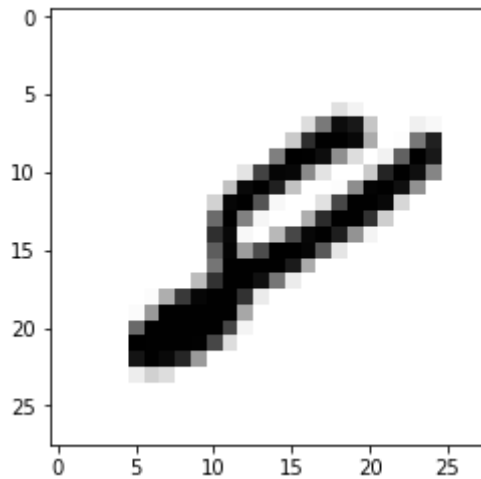
```
In [345]: # Setup train and test splits
```

```
In [346]: (x_train, y_train), (x_test, y_test) = mnist.load_data()
```

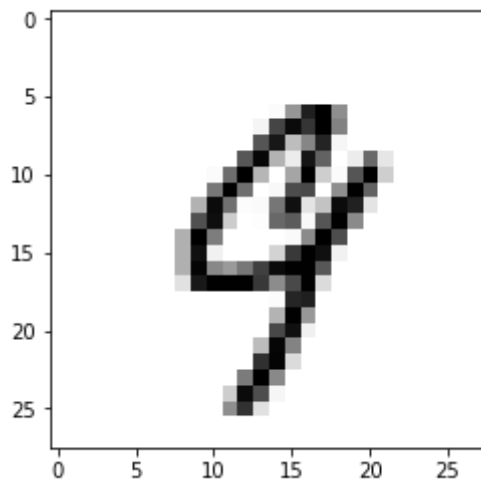
```
In [347]: from random import randint
from matplotlib import pyplot

%matplotlib inline

# protip - visually inspect your data
i = randint(0, x_train.shape[0])
pyplot.imshow(x_train[i], cmap='gray_r');
```



```
In [348]: i = 27_074
pyplot.imshow(x_train[i], cmap='gray_r');
```

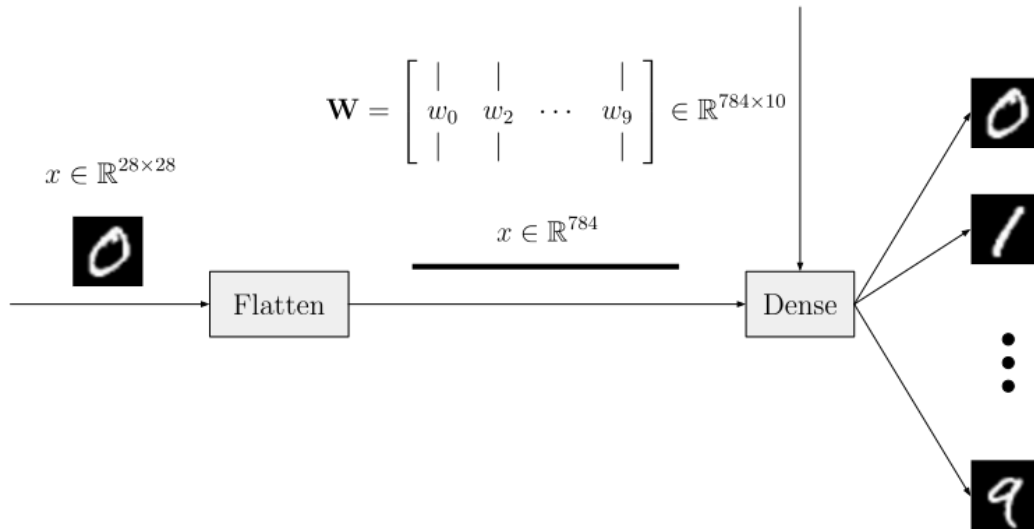


Is this thing a 4 or 9?

```
In [349]: y_train[i]
```

```
Out[349]: 9
```

## Munge data



Convert image matrix into vector to feed into first layer

```
In [350]: # Munge Data
```

```
In [351]: image_size = 784 # 28 x 28
```

```
x_train = x_train.reshape(x_train.shape[0], image_size) # Transform from
matrix to vector
x_train = x_train.astype('float32') # Cast as 32 bit integers
x_train /= 255 # Normalize inputs from 0-255 to 0.0-1.0

x_test = x_test.reshape(x_test.shape[0], image_size) # Transform from ma
trix to vector
x_test = x_test.astype('float32') # Cast as 32 bit integers
x_test /= 255 # Normalize inputs from 0-255 to 0.0-1.0
```

```
In [352]: # Convert class vectors to binary class matrices
```

```
In [353]: y_train = keras.utils.to_categorical(y_train, 10)
y_test = keras.utils.to_categorical(y_test, 10)
```

```
In [354]: # Import the most common type of neural network
```

```
In [355]: from keras.models import Sequential
```

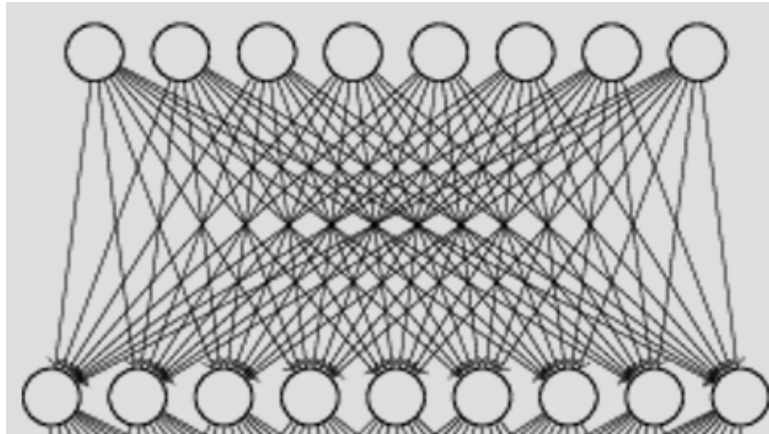
```
In [356]: # Define model instance
```

```
In [357]: model = Sequential()
```

```
In [358]: # Import the most common type of network layer, fully interconnected
```

```
In [359]: from keras.layers import Dense
```

RTFM - <https://keras.io/layers/> (<https://keras.io/layers/>)



```
In [360]: # Define input layer
```

```
In [361]: layer_input = Dense(units=512,                # Number of nodes
                              activation='sigmoid',    # The nonlinearity
                              input_shape=(image_size,))
model.add(layer_input)
```

```
In [362]: # Define another layer
```

```
In [363]: model.add(Dense(units=512, activation='sigmoid'))
```

```
In [364]: # Define output layers
```

```
In [365]: layer_output = Dense(units=10,                # Number of digits (0-9)
                                activation='softmax')   # Convert neural activation t
                                                         o probability of category
model.add(layer_output)
```

```
In [366]: # Print summary of model architecture
```

```
In [367]: model.summary()
```

Layer (type)	Output Shape	Param #
dense_20 (Dense)	(None, 512)	401920
dense_21 (Dense)	(None, 512)	262656
dense_22 (Dense)	(None, 10)	5130

Total params: 669,706  
Trainable params: 669,706  
Non-trainable params: 0

```
In [368]: # Add training parameters to architecture
```

```
In [369]: # Yes - we compile the model to run it
model.compile(loss='categorical_crossentropy',
              optimizer='sgd',
              metrics=['accuracy'])
```

```
In [370]: # Train the model to learn weights
```

```
In [371]: training = model.fit(x_train,
                              y_train,
                              epochs=5, # Number of passes over complete dataset
                              verbose=True,
                              validation_split=0.1)
```

Train on 54000 samples, validate on 6000 samples

Epoch 1/5

54000/54000 [=====] - 17s 312us/step - loss: 2.1620 - acc: 0.3160 - val\_loss: 1.9285 - val\_acc: 0.4772

Epoch 2/5

54000/54000 [=====] - 15s 272us/step - loss: 1.5386 - acc: 0.6498 - val\_loss: 1.1031 - val\_acc: 0.7738

Epoch 3/5

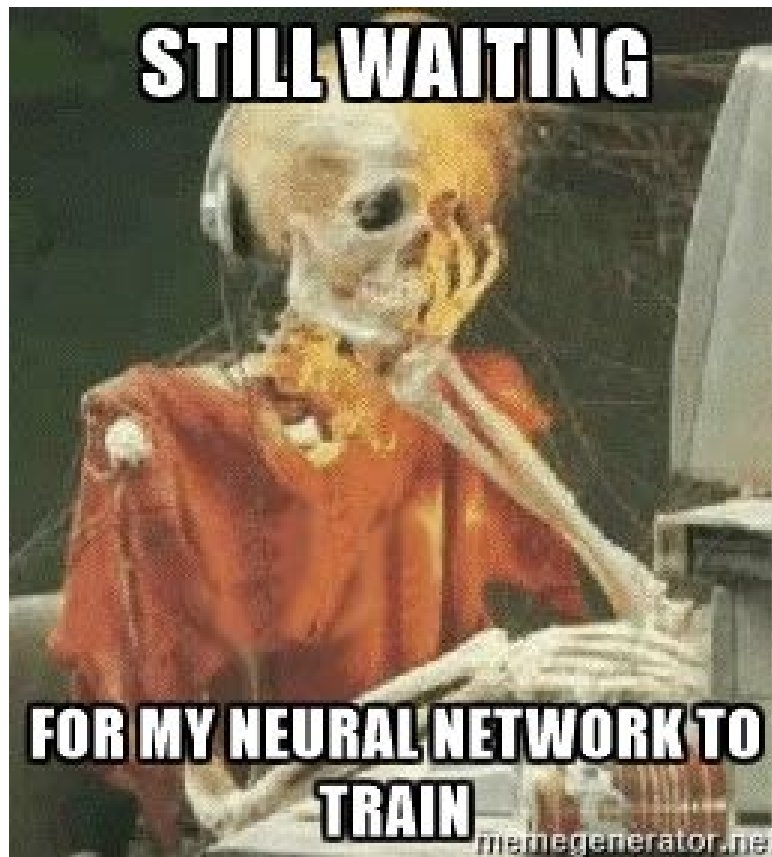
54000/54000 [=====] - 13s 245us/step - loss: 0.9209 - acc: 0.7743 - val\_loss: 0.6921 - val\_acc: 0.8255

Epoch 4/5

54000/54000 [=====] - 15s 284us/step - loss: 0.6692 - acc: 0.8217 - val\_loss: 0.5246 - val\_acc: 0.8690

Epoch 5/5

54000/54000 [=====] - 15s 283us/step - loss: 0.5523 - acc: 0.8482 - val\_loss: 0.4430 - val\_acc: 0.8845



```
In [372]: # Let's see how well our model performs
```

```
In [373]: loss, accuracy = model.evaluate(x_test,
                                           y_test,
                                           verbose=True)

print(f"Test loss: {loss:.3}")
print(f"Test accuracy: {accuracy:.3%}")
```

```
10000/10000 [=====] - 1s 79us/step
Test loss: 0.494
Test accuracy: 86.230%
```

## Keras' Other Features

- Common built-in functions (e.g., activation functions and optimizers)
- Convolutional neural network (CNN or ConvNet)
- Recurrent neural network (RNN) & Long-short term memory (LSTM)
- Pre-trained models



## Summary

- Keras is designed for human beings, not computers.
- Easier to try out Deep Learning (focus on the what, not the how).
- Simple to define neural networks.



**Dmitriy Ryaboy**  
@squarecog

 **Follow**

Replying to @squarecog @josh\_wills

(DL is cool as hell for the right problem. But the number of people who claim knowledge having written 10 lines of Keras...)

RETWEET

1

LIKES

8



9:45 PM - 17 Feb 2017



1



1



8

## Futher Study - Keras

- Keras docs (<https://keras.io/>)
- Keras blog (<https://blog.keras.io/>)
- Keras courses
  - edX (<https://www.edx.org/course/deep-learning-fundamentals-with-keras>)
  - Coursera (<https://www.coursera.org/lecture/ai/keras-overview-7GfN9>)

## Futher Study - Deep Learning

- Prerequisites: Linear Algebra, Probability, Machine Learning
- fast.ai Course (<http://www.fast.ai/>)
- Deep Learning Book (<http://www.deeplearningbook.org/>)



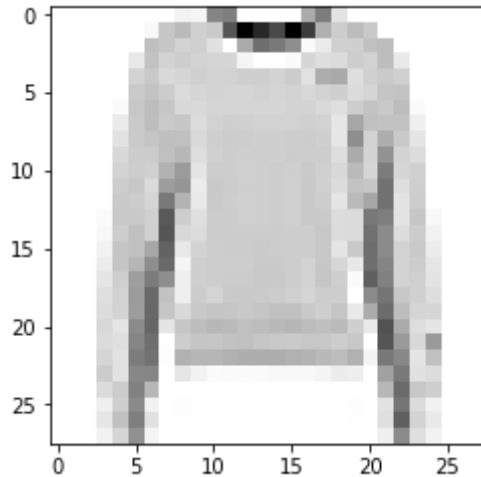
THANK  
YOU  
AND  
ANY  
QUESTIONS

## Bonus Material

```
In [389]: from keras.datasets import fashion_mnist
```

```
In [390]: # Setup train and test splits  
(x_train, y_train), (x_test, y_test) = fashion_mnist.load_data()
```

```
In [391]: pyplot.imshow(x_train[randint(0, x_train.shape[0])], cmap='gray_r');
```



```
In [392]: # Define CNN model

# Redefine input dimensions to make sure conv works
img_rows, img_cols = 28, 28
x_train = x_train.reshape(x_train.shape[0], img_rows, img_cols, 1)
x_test = x_test.reshape(x_test.shape[0], img_rows, img_cols, 1)
input_shape = (img_rows, img_cols, 1)
```

```
In [393]: # Convert class vectors to binary class matrices
y_train = keras.utils.to_categorical(y_train, 10)
y_test = keras.utils.to_categorical(y_test, 10)
```

```
In [394]: from keras.layers import Conv2D, Dense, Flatten, MaxPooling2D
```

```
In [398]: # Define model
model = Sequential()
model.add(Conv2D(32,
                 kernel_size=(3, 3),
                 activation='relu',
                 input_shape=input_shape))
model.add(Conv2D(64, (3, 3), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Flatten())
model.add(Dense(128, activation='relu'))
model.add(Dense(10, activation='softmax'))
```

```
In [401]: model.summary()
```

Layer (type)	Output Shape	Param #
conv2d_5 (Conv2D)	(None, 26, 26, 32)	320
conv2d_6 (Conv2D)	(None, 24, 24, 64)	18496
max_pooling2d_3 (MaxPooling2D)	(None, 12, 12, 64)	0
flatten_3 (Flatten)	(None, 9216)	0
dense_25 (Dense)	(None, 128)	1179776
dense_26 (Dense)	(None, 10)	1290

=====  
Total params: 1,199,882  
Trainable params: 1,199,882  
Non-trainable params: 0  
=====

```
In [402]: model.compile(loss='categorical_crossentropy',  
                        optimizer='adam',  
                        metrics=['accuracy'])
```

```
File "<ipython-input-402-d99bd37e4570>", line 3  
    metrics=['accuracy'])  
    ^
```

SyntaxError: invalid syntax

```
In [403]: # Define training  
training = model.fit(x_train,  
                    y_train,  
                    epochs=1,  
                    verbose=True,  
                    validation_split=0.1)
```

Train on 54000 samples, validate on 6000 samples  
Epoch 1/1  
54000/54000 [=====] - 171s 3ms/step - loss: 0.4744 - acc: 0.8259 - val\_loss: 0.4301 - val\_acc: 0.8398

```
In [404]: loss, accuracy = model.evaluate(x_test,  
                                         y_test,  
                                         verbose=True)  
  
print(f"Test loss: {loss:.3}")  
print(f"Test accuracy: {accuracy:.3%}")
```

10000/10000 [=====] - 8s 826us/step  
Test loss: 0.448  
Test accuracy: 83.590%

## What is keras?



**Keras (κέρας) means horn in Greek.**

**It is a reference to a literary image from ancient Greek and Latin literature.**

**First found in the Odyssey, where dream spirits (Oneiroi, singular Oneiros) are divided between those who deceive men with false visions, who arrive to Earth through a gate of ivory, and those who announce a future that will come to pass, who arrive through a gate of horn.**

**It's a play on the words κέρας (horn) / κραίνω (fulfill), and έλέφας (ivory) / έλεφαίρομαι (deceive).**

**Source (<https://keras.io/#why-this-name-keras>)**