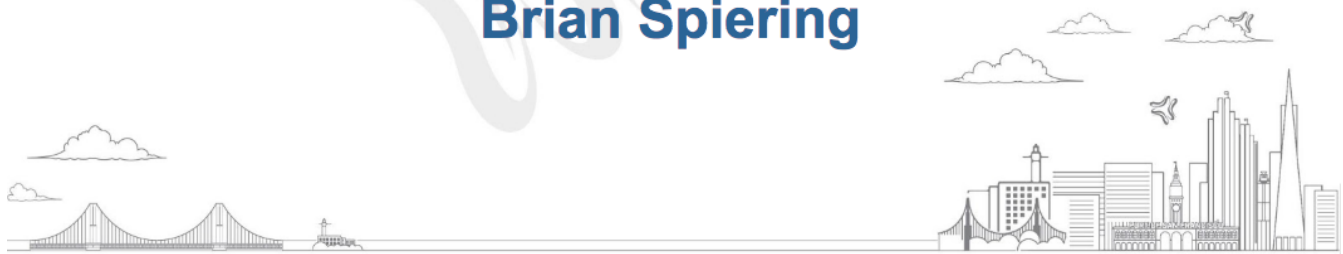




An Absolute Beginner's Guide to Deep Learning with Keras

Brian Spiering



3rd Annual Regional Python Conference

August 16 - 19, 2018 | San Francisco, CA

bit.ly/pybay-keras (<http://bit.ly/pybay-keras>)

Who Am I?

Brian Spiering

What Do I Do?

Professor @



Deep Learning



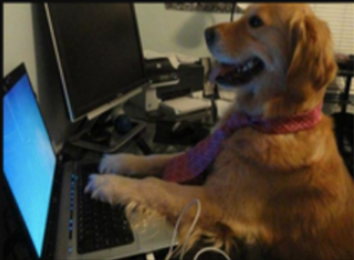
What society thinks I do



What my friends think I do



What other computer scientists think I do



What mathematicians think I do



What I think I do

```
from keras import *
```

What I actually do

Keras - Neural Networks for humans



A high-level, intuitive API for Deep Learning.

Easy to define neural networks, then automatically handles execution.

A simple, modular interface which allows focus on learning and enables fast experimentation

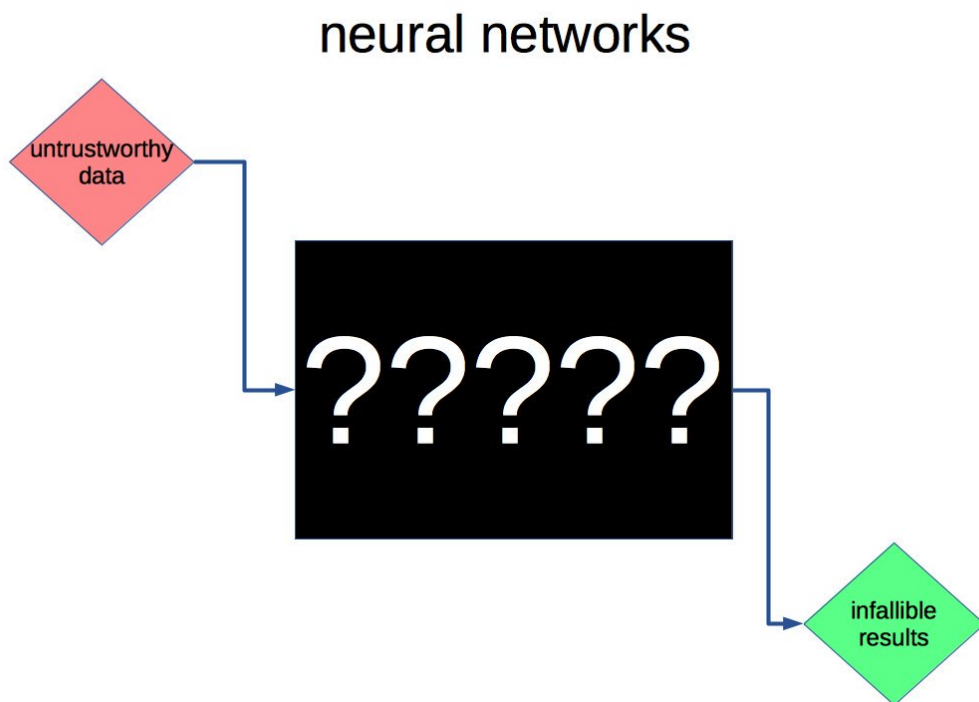
Goals

- General introduction to Deep Learning
- Overview of keras library
- An end-to-end example in keras

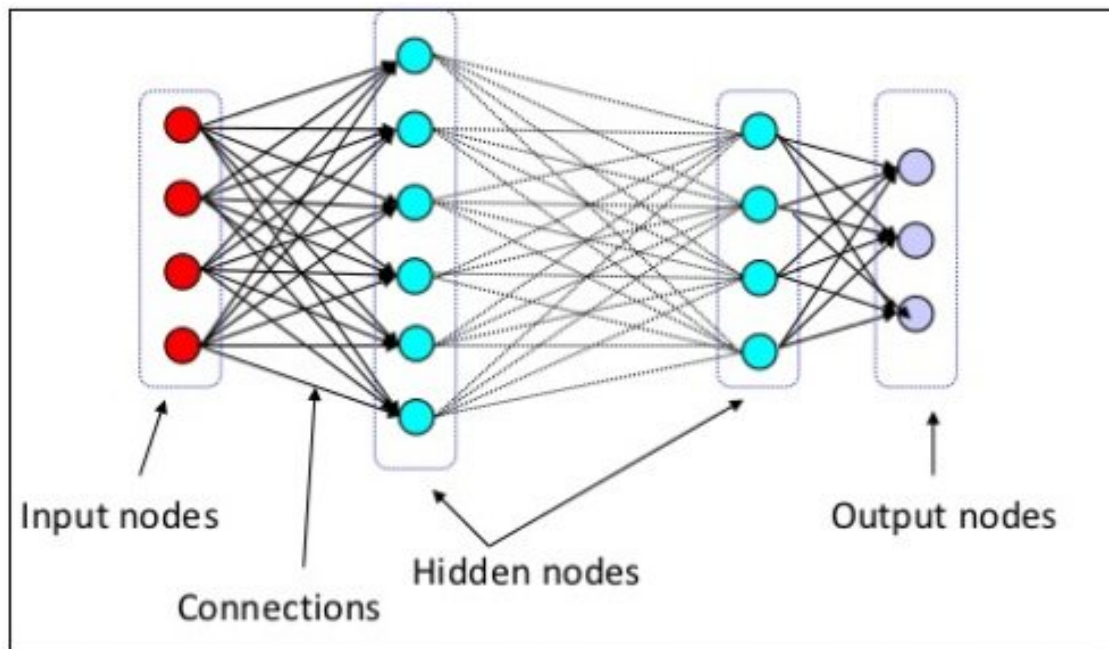
Anti-Goals

- Understanding of Deep Learning (there will be no equations)
- Building neural networks from scratch
- Complete survey of keras library

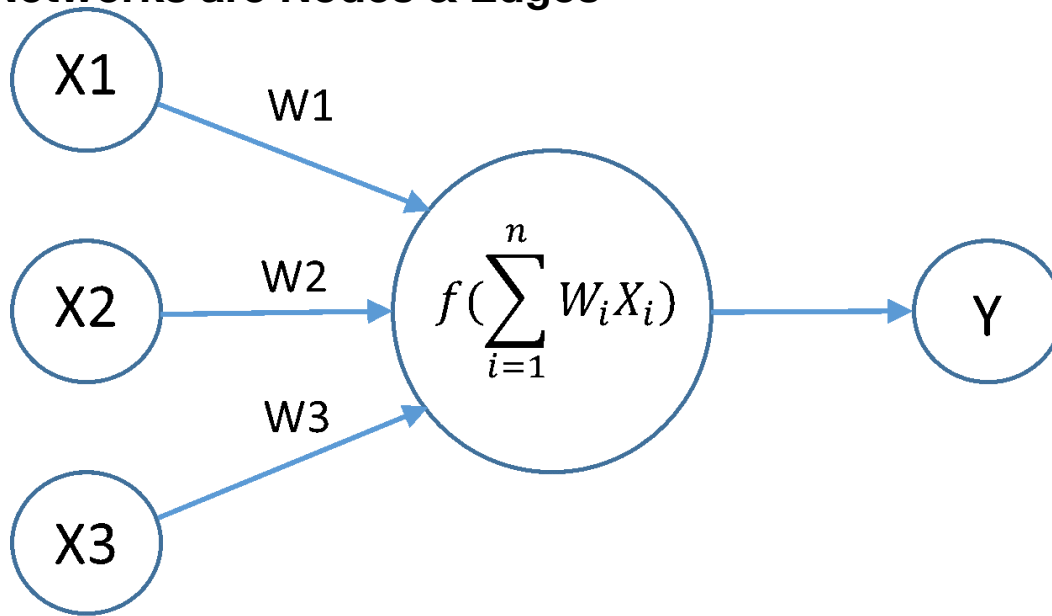
Deep Learning 101



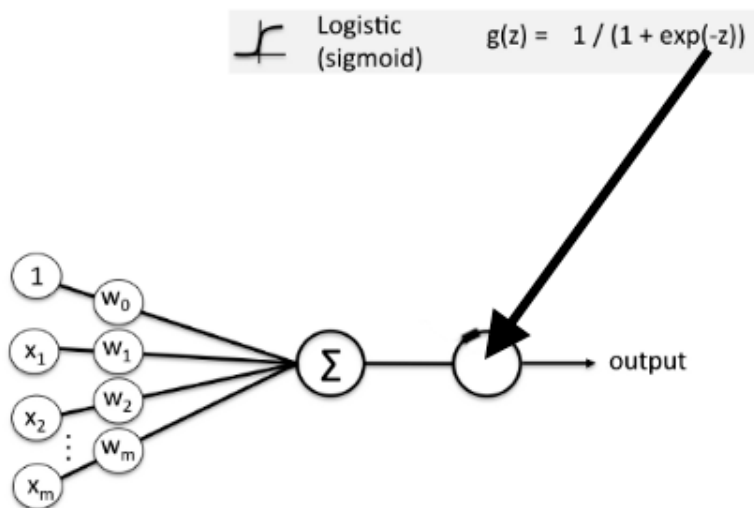
Deep Learning (DL) are Neural networks (NN) with >1 hidden layer



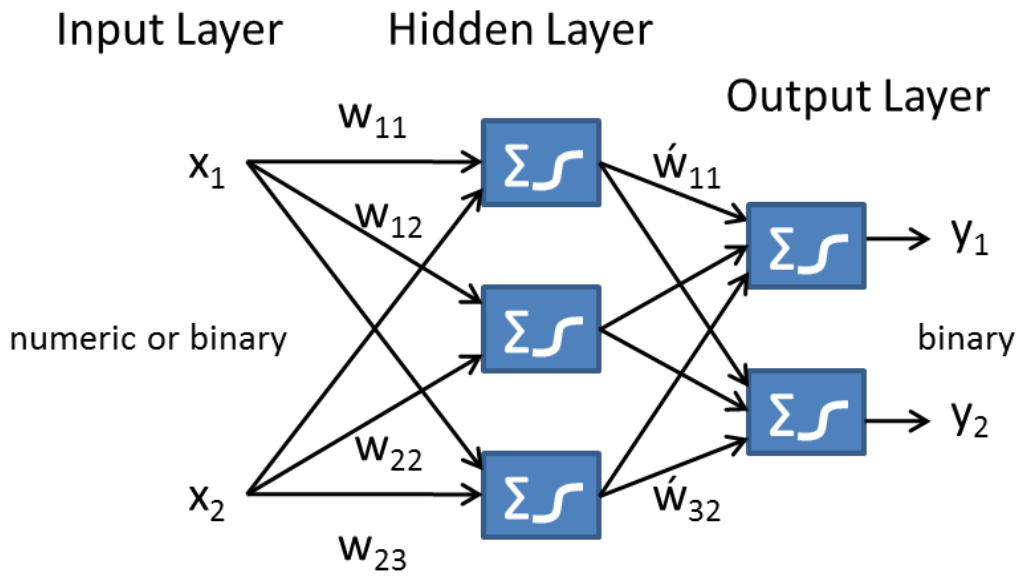
Neural Networks are Nodes & Edges



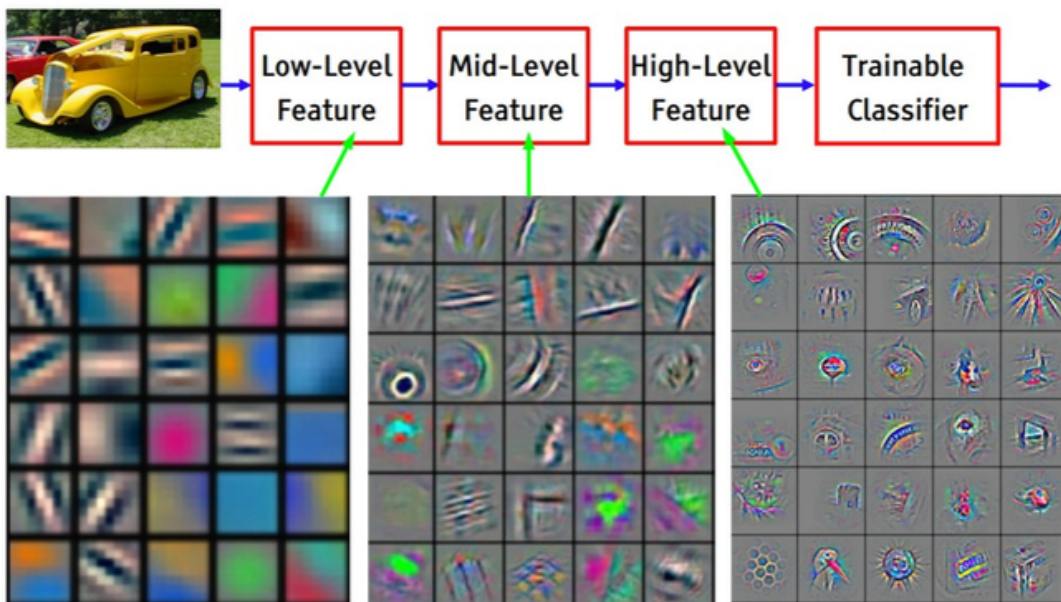
Nonlinear function allows learning of nonlinear relationships



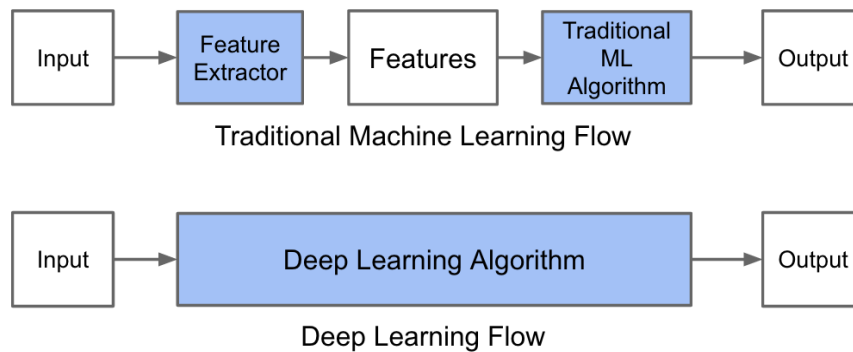
Groups of nodes all the way down



Deep Learning isn't magic, it is just very good at finding patterns



Deep Learning has fewer steps than traditional Machine Learning



If you want to follow along...

GitHub repo: bit.ly/pybay-keras (<http://bit.ly/pybay-keras>)

If you want to type along...

1. Run a local Jupyter Notebook
2. [Binder \(https://mybinder.org/v2/gh/brianspiering/keras-intro/master\)](https://mybinder.org/v2/gh/brianspiering/keras-intro/master): In-Browser Jupyter Notebook
3. [Colaboratory \(https://colab.research.google.com/\)](https://colab.research.google.com/): "Google Docs for Jupyter Notebooks"

```
In [84]: reset -fs
```

```
In [85]: import keras
```

```
In [86]: # What is the backend / execution engine?
```

```
In [87]: keras.backend.backend()
```

```
Out[87]: 'tensorflow'
```



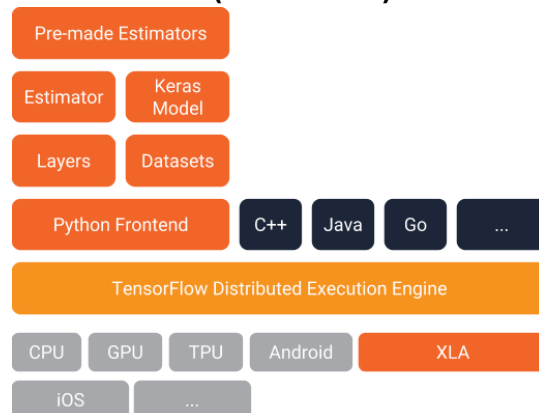
"An open-source software library for Machine Intelligence"

Numerical computation using data flow graphs.

TensorFlow: A great backend

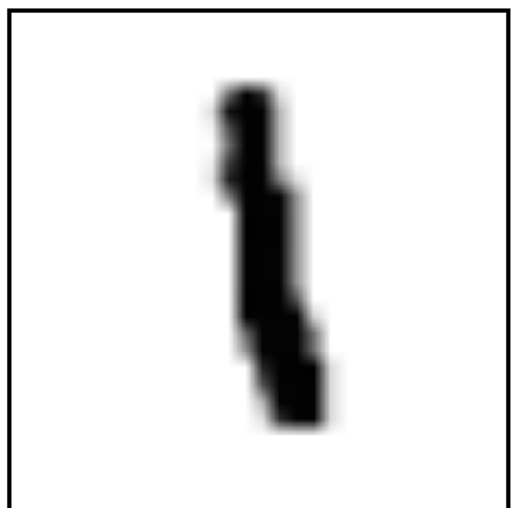
A very flexible architecture which allows you to do almost any numerical operation.

Then deploy the computation to CPUs or GPUs (one or more) across desktop, cloud, or mobile device.



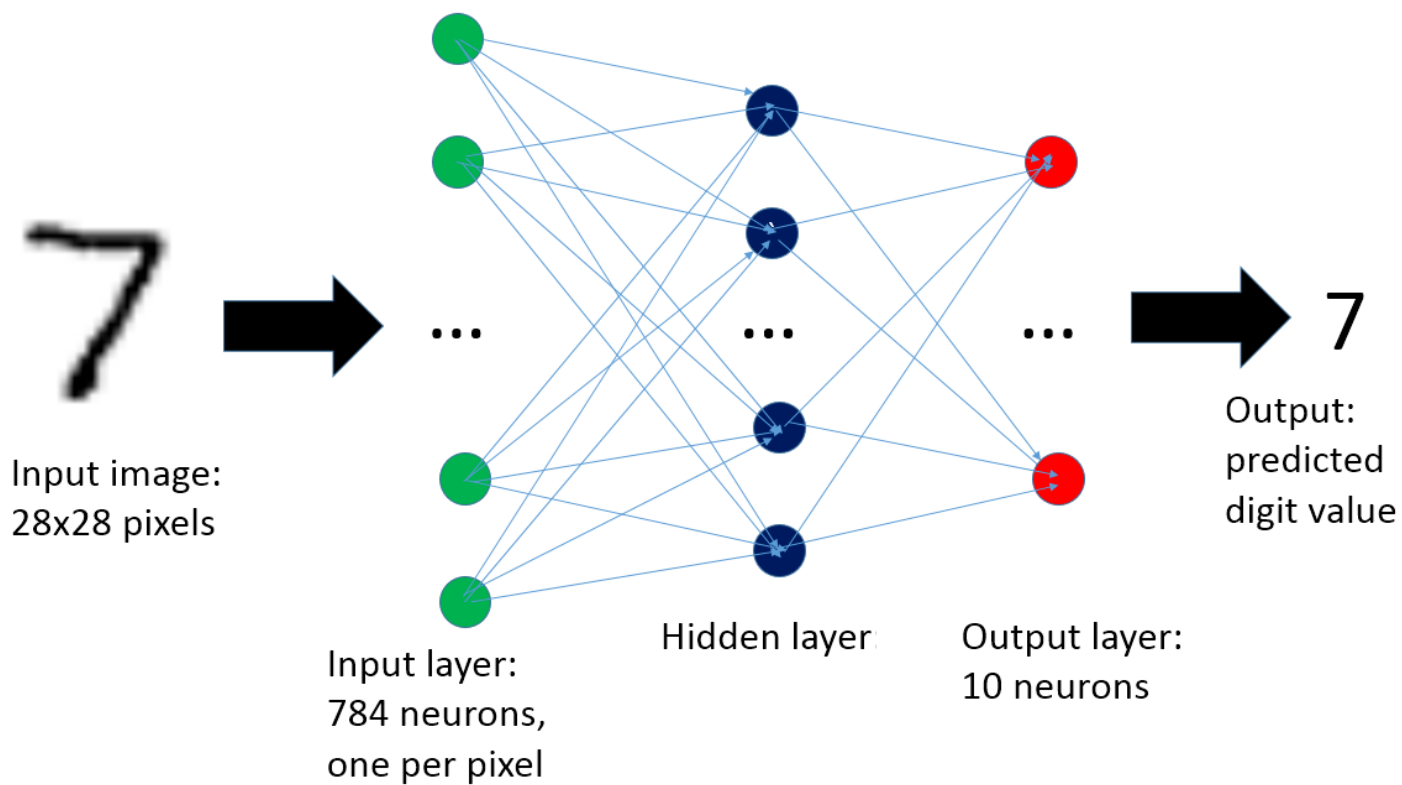
MNIST handwritten digit database:
The "Hello World!" of Computer Vision

2	1	0	4	1	4	9	5
9	0	6	9	0	1	5	9
7	3	4	9	6	6	5	4
0	7	4	0	1	3	1	3
4	7	2	7	1	2	1	1
7	4	2	3	5	1	2	4
4	6	3	5	5	6	0	4
1	9	5	7	8	9	3	7



21

[illegible]



```
In [88]: # Import data
```

```
In [89]: from keras.datasets import mnist
```

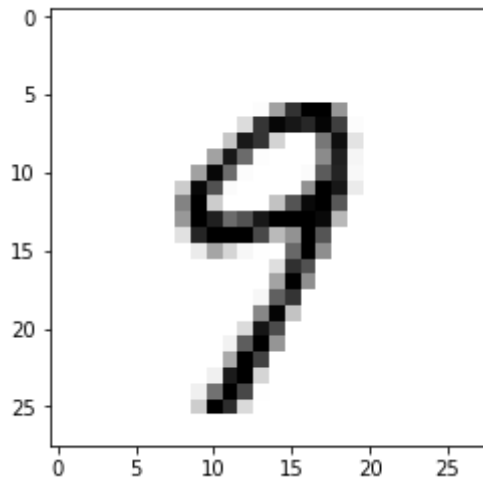
```
In [90]: # Setup train and test splits
```

```
In [91]: (x_train, y_train), (x_test, y_test) = mnist.load_data()
```

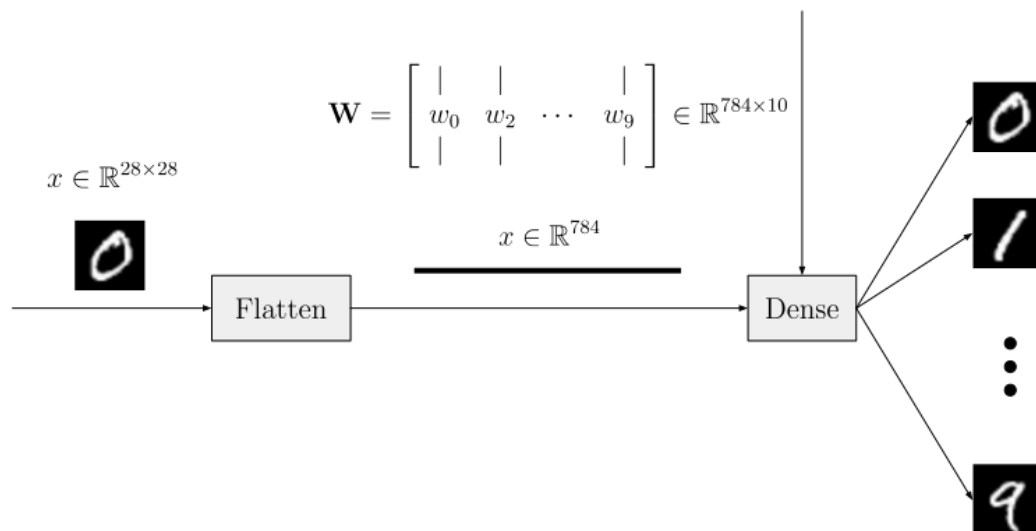
```
In [92]: from random import randint
from matplotlib import pyplot

%matplotlib inline
```

```
In [93]: pyplot.imshow(x_train[randint(0, x_train.shape[0])], cmap='gray_r');
```



Munge data



Convert image matrix into vector to feed into first layer

```
In [94]: # Munge Data
# Transform from matrix to vector, cast, and normalize
```

```
In [95]: image_size = 784 # 28 x 28

x_train = x_train.reshape(x_train.shape[0], image_size) # Transform from
matrix to vector
x_train = x_train.astype('float32') # Cast as 32 bit integers
x_train /= 255 # Normalize inputs from 0-255 to 0.0-1.0

x_test = x_test.reshape(x_test.shape[0], image_size) # Transform from ma
trix to vector
x_test = x_test.astype('float32') # Cast as 32 bit integers
x_test /= 255 # Normalize inputs from 0-255 to 0.0-1.0
```

```
In [96]: # Convert class vectors to binary class matrices
```

```
In [97]: y_train = keras.utils.to_categorical(y_train, 10)
y_test = keras.utils.to_categorical(y_test, 10)
```

```
In [98]: # Import the most common type of neural network
```

```
In [99]: from keras.models import Sequential
```

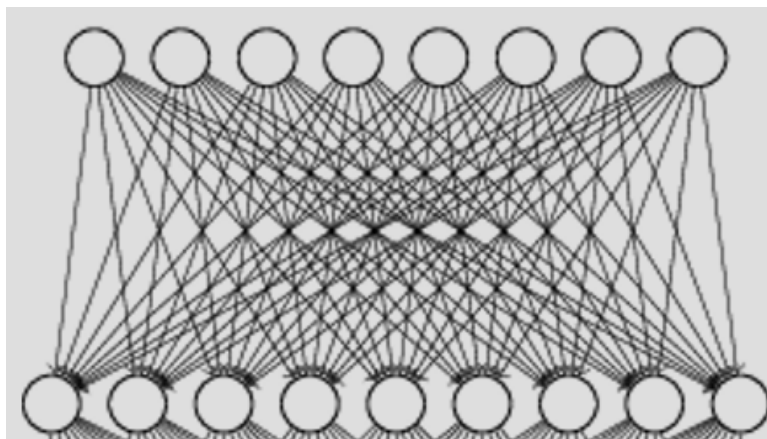
RTFM - <https://keras.io/layers/> (<https://keras.io/layers/>)

```
In [100]: # Define model instance
```

```
In [101]: model = Sequential()
```

```
In [102]: # Import the most common type of network layer, fully interconnected
```

```
In [103]: from keras.layers import Dense
```



```
In [104]: # Define input layer
```

```
In [105]: layer_input = Dense(units=512,          # Number of nodes
                             activation='sigmoid', # The nonlinearity
                             input_shape=(image_size,))
model.add(layer_input)
```

```
In [106]: # Define another layer
```

```
In [107]: model.add(Dense(units=512, activation='sigmoid'))
```

```
In [108]: # Define output layers
```

```
In [109]: layer_output = Dense(units=10,          # Number of digits (0-9)
                                activation='softmax') # Convert neural activation to
                                                    # probability of category
model.add(layer_output)
```

```
In [110]: # Print summary
```

```
In [111]: model.summary()
```

Layer (type)	Output Shape	Param #
dense_9 (Dense)	(None, 512)	401920
dense_10 (Dense)	(None, 512)	262656
dense_11 (Dense)	(None, 10)	5130

=====
Total params: 669,706
Trainable params: 669,706
Non-trainable params: 0
=====

```
In [112]: # Yes - we compile the model to run it
```

```
In [113]: model.compile(loss='categorical_crossentropy',
                        optimizer='sgd',
                        metrics=['accuracy'])
```

```
In [114]: # Train the model
```

```
In [115]: training = model.fit(x_train,
                                y_train,
                                epochs=5, # Number of passes over complete dataset
                                verbose=True,
                                validation_split=0.1)
```

Train on 54000 samples, validate on 6000 samples

Epoch 1/5

54000/54000 [=====] - 15s 285us/step - loss: 2.1522 - acc: 0.3213 - val_loss: 1.8987 - val_acc: 0.5315

Epoch 2/5

54000/54000 [=====] - 14s 262us/step - loss: 1.5000 - acc: 0.6548 - val_loss: 1.0769 - val_acc: 0.7430

Epoch 3/5

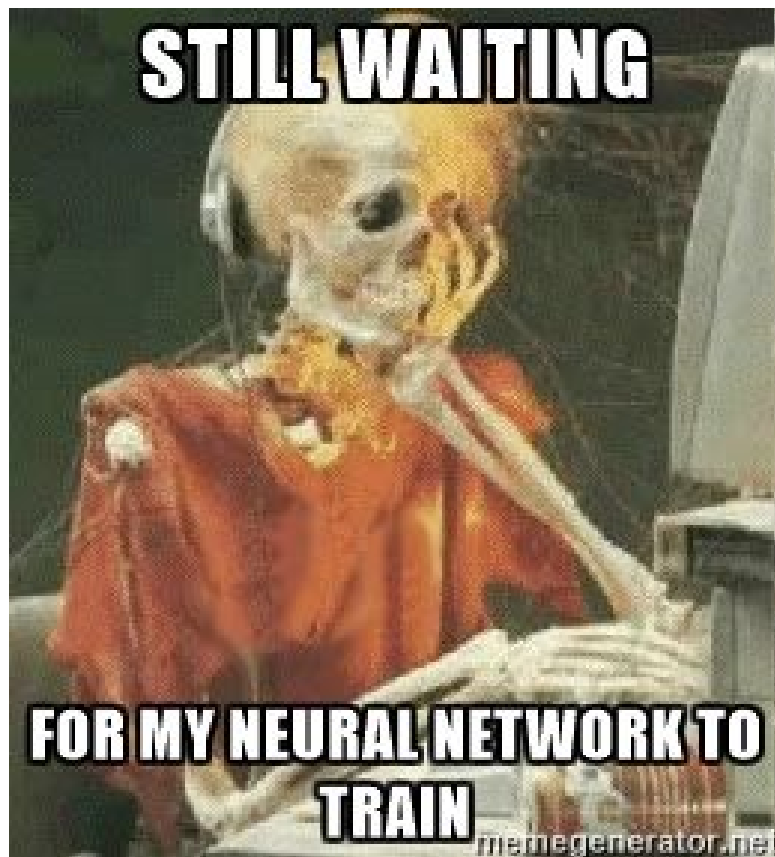
54000/54000 [=====] - 15s 285us/step - loss: 0.9003 - acc: 0.7860 - val_loss: 0.6709 - val_acc: 0.8560

Epoch 4/5

54000/54000 [=====] - 14s 266us/step - loss: 0.6515 - acc: 0.8317 - val_loss: 0.5121 - val_acc: 0.8778

Epoch 5/5

54000/54000 [=====] - 18s 340us/step - loss: 0.5385 - acc: 0.8549 - val_loss: 0.4268 - val_acc: 0.8940



```
In [116]: # Let's see how well our model performs
```

```
In [117]: loss, accuracy = model.evaluate(x_test,
                                           y_test,
                                           verbose=True)

print(f"Test loss: {loss:.3}")
print(f"Test accuracy: {accuracy:.3%}")

10000/10000 [=====] - 1s 106us/step
Test loss: 0.476
Test accuracy: 87.140%
```

Keras' Other Features

- Common built-in functions (e.g., activation functions and optimizers)
- Convolutional neural network (CNN or ConvNet)
- Recurrent neural network (RNN) & Long-short term memory (LSTM)
- Pre-trained models

Summary

- Keras is designed for human beings, not computers.
- Easier to try out Deep Learning (focus on the what, not the how).
- Simple to define neural networks.



Dmitriy Ryaboy
@squarecog

 **Follow**

Replying to @squarecog @josh_wills

(DL is cool as hell for the right problem. But the number of people who claim knowledge having written 10 lines of Keras...)

RETWEET

1

LIKES

8



9:45 PM - 17 Feb 2017



1



1



8

Futher Study - Keras

- [Keras docs \(https://keras.io/\)](https://keras.io/)
- [Keras blog \(https://blog.keras.io/\)](https://blog.keras.io/)
- Keras courses
 - [edX \(https://www.edx.org/course/deep-learning-fundamentals-with-keras\)](https://www.edx.org/course/deep-learning-fundamentals-with-keras)
 - [Coursera \(https://www.coursera.org/lecture/ai/keras-overview-7GfN9\)](https://www.coursera.org/lecture/ai/keras-overview-7GfN9)

Futher Study - Deep Learning

- Prerequisites: Linear Algebra, Probability, Machine Learning
- [fast.ai Course \(http://www.fast.ai/\)](http://www.fast.ai/)
- [Deep Learning Book \(http://www.deeplearningbook.org/\)](http://www.deeplearningbook.org/)

Bonus Material

```
In [118]: # reset -fs
```

```
In [119]: # from keras import *
```

```
In [120]: # whos
```

```
In [121]: # from keras.datasets import fashion_mnist
```

```
In [122]: # # Setup train and test splits
# (x_train, y_train), (x_test, y_test) = fashion_mnist.load_data()
```

```
In [123]: # from random import randint
# from matplotlib import pyplot

# %matplotlib inline
```

```
In [124]: # pyplot.imshow(x_train[randint(0, x_train.shape[0])], cmap='gray_r');
```

```
In [125]: # # Define CNN model

# # Redefine input dimensions to make sure conv works
# img_rows, img_cols = 28, 28
# x_train = x_train.reshape(x_train.shape[0], img_rows, img_cols, 1)
# x_test = x_test.reshape(x_test.shape[0], img_rows, img_cols, 1)
# input_shape = (img_rows, img_cols, 1)
```

```
In [126]: # import keras
```

```
In [127]: # # Convert class vectors to binary class matrices
# y_train = keras.utils.to_categorical(y_train, 10)
# y_test = keras.utils.to_categorical(y_test, 10)
```

```
In [128]: # from keras.layers import Conv2D, Dense, Flatten, MaxPooling2D
```

```
In [129]: # # Define model
# model = Sequential()
# model.add(Conv2D(32,
#                 kernel_size=(3, 3),
#                 activation='sigmoid',
#                 input_shape=input_shape))
# model.add(Conv2D(64, (3, 3), activation='sigmoid'))
# model.add(MaxPooling2D(pool_size=(2, 2)))
# model.add(Flatten())
# model.add(Dense(128, activation='sigmoid'))
# model.add(Dense(10, activation='softmax'))
```

```
In [130]: # model.compile(loss='categorical_crossentropy',
#                       optimizer='adam',
#                       metrics=['accuracy'])
```

```
In [131]: # # Define training
# training = model.fit(x_train,
#                     y_train,
#                     epochs=5,
#                     verbose=True,
#                     validation_split=0.1)
```

```
In [132]: # loss, accuracy = model.evaluate(x_test,
#                                           y_test,
#                                           verbose=True)
# print(f"Test loss: {loss:.3}")
# print(f"Test accuracy: {accuracy:.3%}")
```

What is keras?



Keras (κέρας) means horn in Greek.

It is a reference to a literary image from ancient Greek and Latin literature.

First found in the Odyssey, where dream spirits (Oneiroi, singular Oneiros) are divided between those who deceive men with false visions, who arrive to Earth through a gate of ivory, and those who announce a future that will come to pass, who arrive through a gate of horn.

It's a play on the words κέρας (horn) / κραίνω (fulfill), and έλέφας (ivory) / έλεφαίρομαι (deceive).

Source (<https://keras.io/#why-this-name-keras>)