

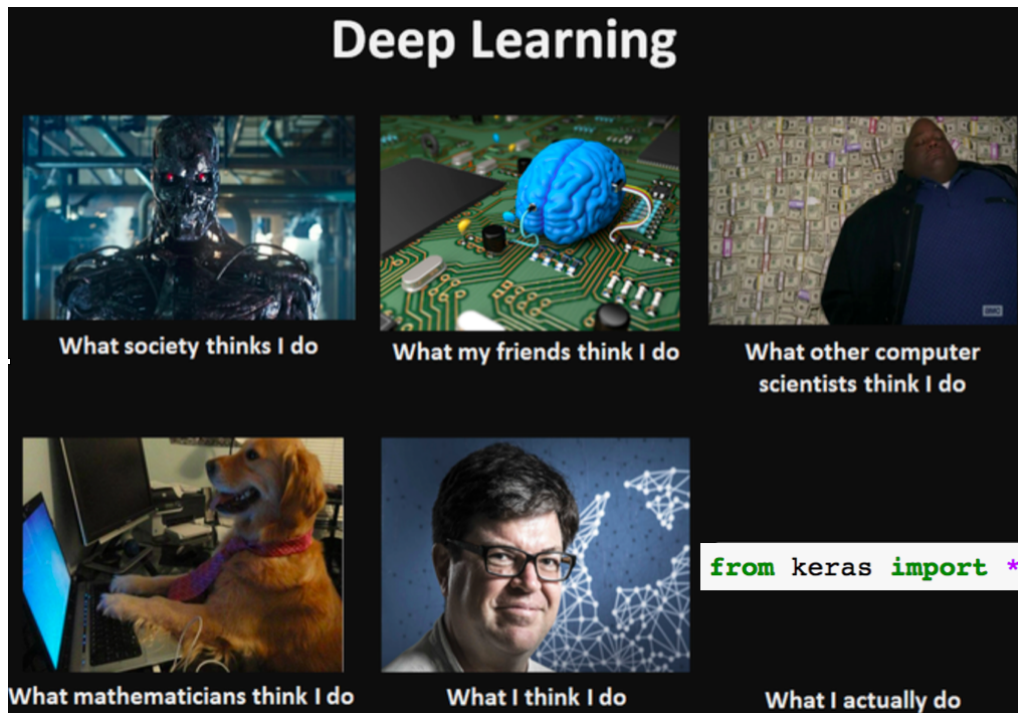
Who Am I?

Brian Spiering


What Do I Do?

Professor @





Keras - Neural Networks for humans




A high-level, intuitive API for Deep Learning.

Easy to define neural networks, then automatically handles execution.


A simple, modular interface which allows focus on learning and enables fast experimentation

Deep Learning 101


neural networks




Deep Learning (DL) are Neural networks (NN) with >1 hidden layer




Neural Networks are Nodes & Edges



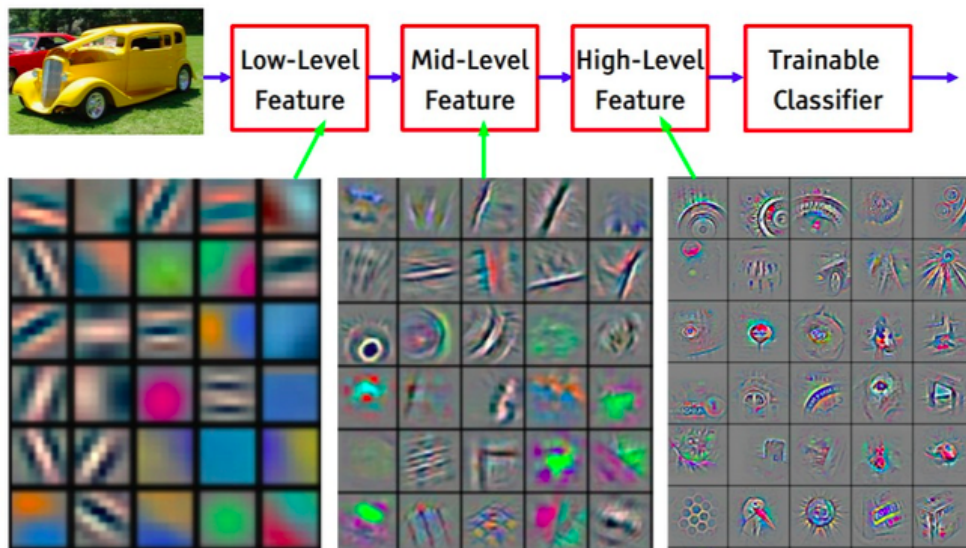
Nonlinear function allows learning of nonlinear relationships



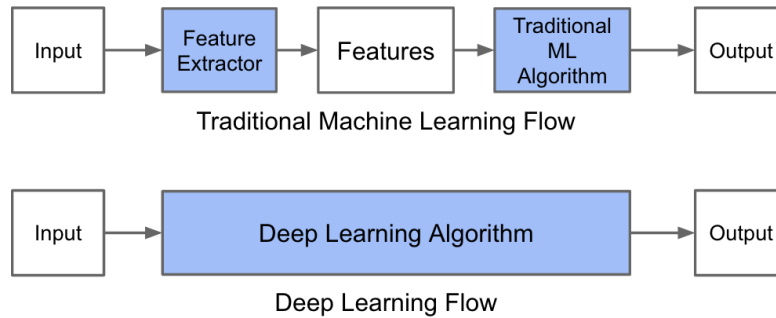
Groups of nodes all the way down



Deep Learning isn't magic, it is just very good at finding patterns



Deep Learning has fewer steps than traditional Machine Learning



If you want to follow along...

GitHub repo: bit.ly/pybay-keras (<http://bit.ly/pybay-keras>)

If you want to type along...

1. Run a local Jupyter Notebook
2. [Binder \(https://mybinder.org/v2/gh/brianspiering/keras-intro/master\)](https://mybinder.org/v2/gh/brianspiering/keras-intro/master): In-Browser Jupyter Notebook
3. [Colaboratory \(https://colab.research.google.com/\)](https://colab.research.google.com/): "Google Docs for Jupyter Notebooks"

```
In [1]: reset -fs
```

```
In [2]: import keras
```

```
/Users/brian/anaconda3/envs/keras/lib/python3.6/importlib/_bootstrap.py:205: Run
timeWarning: numpy.dtype size changed, may indicate binary incompatibility. Expe
cted 96, got 88
    return f(*args, **kwargs)
Using TensorFlow backend.
```

```
In [3]: # What is the backend / execution engine?
```

```
In [4]: keras.backend.backend()
```

```
Out[4]: 'tensorflow'
```



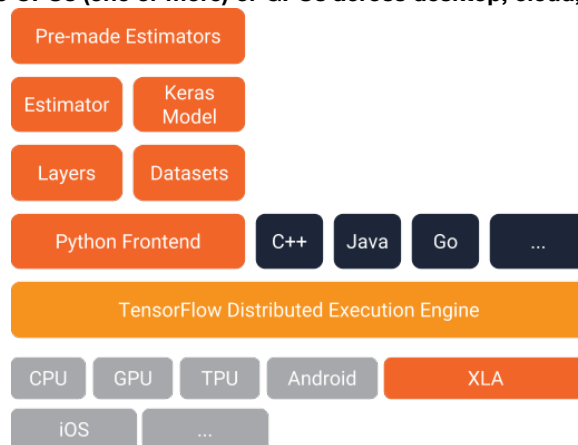
"An open-source software library for Machine Intelligence"

Numerical computation using data flow graphs.

TensorFlow: A great backend

A very flexible architecture that allows you to do almost any operation.


Then deploy the computation to CPUs (one or more) or GPUs across desktop, cloud, or mobile device.



MNIST handwritten digit database:
The “Hello World!” of Computer Vision



$$12 \begin{bmatrix} 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & .6 & .8 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & .7 & 1 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & .7 & 1 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & .5 & 1 & .4 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & .4 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & .4 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & .7 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 1 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & .9 & 1 & .1 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & .3 & 1 & .1 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix}$$



```
In [5]: # Import data
```

```
In [6]: from keras.datasets import mnist
```


```
In [7]: # Setup train and test splits
```

```
In [8]: (x_train, y_train), (x_test, y_test) = mnist.load_data()
```

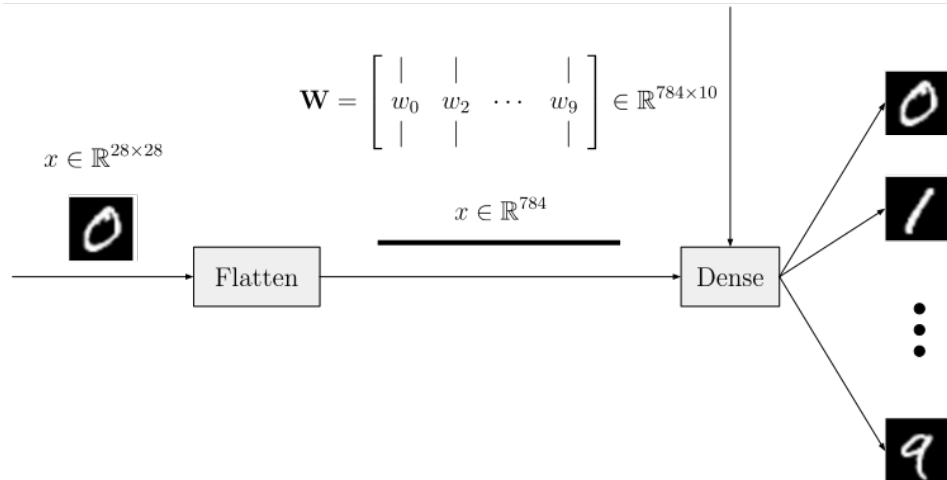
```
In [9]: from random import randint
from matplotlib import pyplot

%matplotlib inline
```

```
In [10]: pyplot.imshow(x_train[randint(0, x_train.shape[0])], cmap='gray_r');
```



Munge data



Convert image matrix into vector to feed into first layer

```
In [11]: # Munge Data
# Transform from matrix to vector, cast, and normalize
```

```
In [12]: image_size = 784 # 28 x 28

x_train = x_train.reshape(x_train.shape[0], image_size) # Transform from matrix to vector
x_train = x_train.astype('float32') # Cast as 32 bit integers
x_train /= 255 # Normalize inputs from 0-255 to 0.0-1.0

x_test = x_test.reshape(x_test.shape[0], image_size) # Transform from matrix to vector
x_test = x_test.astype('float32') # Cast as 32 bit integers
x_test /= 255 # Normalize inputs from 0-255 to 0.0-1.0
```

```
In [13]: # Convert class vectors to binary class matrices
```

```
In [14]: y_train = keras.utils.to_categorical(y_train, 10)
y_test = keras.utils.to_categorical(y_test, 10)
```

```
In [15]: # Import the most common type of neural network
```

```
In [16]: from keras.models import Sequential
```


RTFM - <https://keras.io/layers/> (<https://keras.io/layers/>)

```
In [17]: # Define model instance
```

```
In [18]: model = Sequential()
```

```
In [19]: # Import the most common type of network layer, fully interconnected
```

```
In [20]: from keras.layers import Dense
```



```
In [21]: # Define input layer
```

```
In [22]: layer_input = Dense(units=512,          # Number of nodes
                             activation='sigmoid', # The nonlinearity
                             input_shape=(image_size,))
model.add(layer_input)
```

```
In [23]: # Define another layer
```

```
In [24]: model.add(Dense(units=512, activation='sigmoid'))
```

```
In [25]: # Define output layers
```

```
In [26]: layer_output = Dense(units=10,          # Number of digits (0-9)
                               activation='softmax') # Convert neural activation to probability of category
model.add(layer_output)
```

```
In [27]: # Print summary
```

```
In [28]: model.summary()
```

Layer (type)	Output Shape	Param #
dense_1 (Dense)	(None, 512)	401920
dense_2 (Dense)	(None, 512)	262656
dense_3 (Dense)	(None, 10)	5130
Total params: 669,706		
Trainable params: 669,706		
Non-trainable params: 0		

```
In [29]: # Yes - we compile the model to run it
```

```
In [30]: model.compile(loss='categorical_crossentropy',
                      optimizer='sgd',
                      metrics=['accuracy'])
```

```
In [31]: # Train the model
```

```
In [32]: training = model.fit(x_train,  
                               y_train,  
                               epochs=5, # Number of passes over complete dataset  
                               verbose=True,  
                               validation_split=0.1)
```

Train on 54000 samples, validate on 6000 samples

Epoch 1/5

54000/54000 [=====] - 13s 236us/step - loss: 2.1548 - acc: 0.3199 - val_loss: 1.9109 - val_acc: 0.4287

Epoch 2/5

54000/54000 [=====] - 13s 247us/step - loss: 1.5084 - acc: 0.6597 - val_loss: 1.0729 - val_acc: 0.7858

Epoch 3/5


54000/54000 [=====] - 14s 250us/step - loss: 0.9052 - acc: 0.7871 - val_loss: 0.6801 - val_acc: 0.8432

Epoch 4/5

54000/54000 [=====] - 14s 259us/step - loss: 0.6555 - acc: 0.8338 - val_loss: 0.5119 - val_acc: 0.8820

Epoch 5/5

54000/54000 [=====] - 13s 248us/step - loss: 0.5399 - acc: 0.8579 - val_loss: 0.4344 - val_acc: 0.8862



```
In [33]: # Let's see how well our model performs
```

```
In [34]: loss, accuracy = model.evaluate(x_test,  
                                         y_test,  
                                         verbose=True)  
  
print(f"Test loss: {loss:.3}")  
print(f"Test accuracy: {accuracy:.3%}")
```

10000/10000 [=====] - 1s 73us/step

Test loss: 0.482

Test accuracy: 86.530%

Keras' Other Features

- Common built-in functions (e.g., activation functions and optimizers)
- Convolutional neural network (CNN or ConvNet)
- Recurrent neural network (RNN) & Long-short term memory (LSTM)
- Pre-trained models

Summary

- Keras is designed for human beings, not computers.
- Easier to try out Deep Learning (focus on the what, not the how).
- Simple to define neural networks.



Dmitriy Ryaboy
@squarecog

Follow

Replying to @squarecog @josh_wills

(DL is cool as hell for the right problem. But the number of people who claim knowledge having written 10 lines of Keras...)

RETWEET

1

LIKES

8



9:45 PM - 17 Feb 2017



1



1



8

Futher Study - Keras

- Keras docs
 - <https://keras.io/> (<https://keras.io/>)
 - <https://blog.keras.io/> (<https://blog.keras.io/>)
- Keras courses
 - <https://www.edx.org/course/deep-learning-fundamentals-with-keras> (<https://www.edx.org/course/deep-learning-fundamentals-with-keras>)
 - <https://www.coursera.org/lecture/ai/keras-overview-7GfN9> (<https://www.coursera.org/lecture/ai/keras-overview-7GfN9>)

Futher Study - Deep Learning

- Linear Algebra, Probability, Machine Learning
- Study Deep Learning
 - [fast.ai Course](http://www.fast.ai/) (<http://www.fast.ai/>)
 - [Deep Learning Book](http://www.deeplearningbook.org/) (<http://www.deeplearningbook.org/>)

Bonus Material

```
In [35]: reset -fs
```

```
In [36]: from keras import *
```

```
In [37]: whos
```

Variable	Type	Data/Info
Input	function	<function Input at 0x127447488>
Model	type	<class 'keras.engine.training.Model'>
Sequential	type	<class 'keras.engine.sequential.Sequential'>
absolute_import	_Feature	_Feature((2, 5, 0, 'alpha<...>0, 0, 'alpha', 0), 16384)
activations	module	<module 'keras.activation<...>es/keras/activations.py'>
applications	module	<module 'keras.applicatio<...>pplications/__init__.py'>
backend	module	<module 'keras.backend' f<...>ras/backend/__init__.py'>
callbacks	module	<module 'keras.callbacks' <...>ages/keras/callbacks.py'>
constraints	module	<module 'keras.constraint<...>es/keras/constraints.py'>
datasets	module	<module 'keras.datasets' <...>as/datasets/__init__.py'>
engine	module	<module 'keras.engine' fr<...>eras/engine/__init__.py'>
initializers	module	<module 'keras.initialize<...>s/keras/initializers.py'>
layers	module	<module 'keras.layers' fr<...>eras/layers/__init__.py'>
legacy	module	<module 'keras.legacy' fr<...>eras/legacy/__init__.py'>
losses	module	<module 'keras.losses' fr<...>ackages/keras/losses.py'>
metrics	module	<module 'keras.metrics' f<...>ckages/keras/metrics.py'>
models	module	<module 'keras.models' fr<...>ackages/keras/models.py'>
optimizers	module	<module 'keras.optimizers<...>ges/keras/optimizers.py'>
preprocessing	module	<module 'keras.preprocess<...>eprocessing/__init__.py'>
regularizers	module	<module 'keras.regularize<...>s/keras/regularizers.py'>
utils	module	<module 'keras.utils' fro<...>keras/utils/__init__.py'>
wrappers	module	<module 'keras.wrappers' <...>as/wrappers/__init__.py'>


```
In [38]: from keras.datasets import fashion_mnist
```

```
In [39]: # Setup train and test splits
(x_train, y_train), (x_test, y_test) = fashion_mnist.load_data()
```

```
In [40]: from random import randint
         from matplotlib import pyplot

         %matplotlib inline
```

```
In [41]: pyplot.imshow(x_train[randint(0, x_train.shape[0])], cmap='gray_r');
```



```
In [42]: # Define CNN model

         # Redefine input dimensions to make sure conv works
         img_rows, img_cols = 28, 28
         x_train = x_train.reshape(x_train.shape[0], img_rows, img_cols, 1)
         x_test = x_test.reshape(x_test.shape[0], img_rows, img_cols, 1)
         input_shape = (img_rows, img_cols, 1)
```

```
In [43]: import keras
```

```
In [44]: # Convert class vectors to binary class matrices
         y_train = keras.utils.to_categorical(y_train, 10)
         y_test = keras.utils.to_categorical(y_test, 10)
```

```
In [45]: from keras.layers import Conv2D, Dense, Flatten, MaxPooling2D
```

```
In [46]: # Define model
         model = Sequential()
         model.add(Conv2D(32,
                           kernel_size=(3, 3),
                           activation='sigmoid',
                           input_shape=input_shape))
         model.add(Conv2D(64, (3, 3), activation='sigmoid'))
         model.add(MaxPooling2D(pool_size=(2, 2)))
         model.add(Flatten())
         model.add(Dense(128, activation='sigmoid'))
         model.add(Dense(10, activation='softmax'))
```

```
In [47]: model.compile(loss='categorical_crossentropy',
                       optimizer='adam',
                       metrics=['accuracy'])
```

```
In [48]: # Define training
training = model.fit(x_train,
                    y_train,
                    epochs=5,
                    verbose=True,
                    validation_split=0.1)

Train on 54000 samples, validate on 6000 samples
Epoch 1/5
54000/54000 [=====] - 173s 3ms/step - loss: 2.3262 - ac
c: 0.0986 - val_loss: 2.3071 - val_acc: 0.1055
Epoch 2/5
54000/54000 [=====] - 173s 3ms/step - loss: 2.3100 - ac
c: 0.0996 - val_loss: 2.3076 - val_acc: 0.1003
Epoch 3/5
54000/54000 [=====] - 173s 3ms/step - loss: 2.3107 - ac
c: 0.0994 - val_loss: 2.3047 - val_acc: 0.1032
Epoch 4/5
54000/54000 [=====] - 187s 3ms/step - loss: 2.3111 - ac
c: 0.0989 - val_loss: 2.3290 - val_acc: 0.1055
Epoch 5/5
54000/54000 [=====] - 182s 3ms/step - loss: 2.3101 - ac
c: 0.1000 - val_loss: 2.3131 - val_acc: 0.0925
```

```
In [49]: loss, accuracy = model.evaluate(x_test,
                                         y_test,
                                         verbose=True)

print(f"Test loss: {loss:.3}")
print(f"Test accuracy: {accuracy:.3%}")

10000/10000 [=====] - 9s 936us/step
Test loss: 2.31
Test accuracy: 10.000%
```

What is keras?



Keras (κέρας) means horn in Greek.

It is a reference to a literary image from ancient Greek and Latin literature.

First found in the Odyssey, where dream spirits (Oneiroi, singular Oneiros) are divided between those who deceive men with false visions, who arrive to Earth through a gate of ivory, and those who announce a future that will come to pass, who arrive through a gate of horn.

It's a play on the words κέρασ (horn) / κραινω (fulfill), and έλέφας (ivory) / έλεφαίρομαι (deceive).

Source (<https://keras.io/#why-this-name-keras>)